

SpringOne Platform

Designing, Implementing, and Using Reactive APIs

Ben Hale (@nebhale)
Paul Harris (@twoseat)

Pivotal

Introduction to Reactive Programming

- Reactive is used to broadly define event-driven systems
- Moves imperative logic to
 - asynchronous
 - non-blocking
 - functional-style code
- Allows stable, scalable access to external systems

Cloud Foundry Java Client

- Cloud Foundry is composed of
 - RESTful APIs
 - Secured by OAuth
 - Transmitting large JSON request and response payloads
- APIs are low level, often requiring orchestration

Cloud Foundry Java Client

- Java has become the preferred language for orchestrating Cloud Foundry
- Complex orchestration is difficult when working against generic payloads
- A strongly typed Java API is required for maintainability

Generic Request and Response Payloads

```
Map<String, Object> request = new HashMap<>();  
request.put("name", "test-application");  
request.put("application", Paths.get("target/test-application.jar"));  
request.put("buildpack", "java-buildpack");
```

```
this.operations.applications()  
    .create(request)  
    .map(response -> response.get("id"));
```

Strongly Typed Request and Response Payloads

```
this.operations.applications()  
    .create(CreateApplicationRequest.builder()  
        .name("test-application")  
        .application(Paths.get("target/test-application.jar"))  
        .buildpack("java-buildpack")  
        .build())  
    .map(CreateApplicationResponse::getId);
```

Cloud Foundry Java Client

- The high-latency network interaction is a good fit for a Reactive Design
- Encourages asynchronous and highly concurrent usage
- Partnering with Project Reactor early influenced both the design of the client and of Reactor

SpringOne Platform

Design

Pivotal

When To Use Reactive

- Networking
 - Latency
 - Failures
 - Backpressure
- Highly Concurrent Operations
- Scalability

What Should a Reactive API Return?

- `Flux<T>`
 - Returns 0 to N values
- `Mono<T>`
 - Returns 0 to 1 value
- The Java Client often returns Monos containing collections

Using a Flux<T> Response

```
Flux<Application> listApplications() {...}
```

```
Flux<String> listApplicationNames() {  
    return listApplications()  
        .map(Application::getName);  
}
```

```
void printApplicationName() {  
    listApplicationNames()  
        .subscribe(System.out::println);  
}
```

Using a Mono<T> Response

```
Flux<Application> listApplications() {...}
```

```
Mono<List<String>> listApplicationNames() {  
    return listApplications()  
        .map(Application::getName)  
        .collectAsList();  
}
```

```
Mono<Boolean> doesApplicationExist(String name) {  
    return listApplicationNames()  
        .map(names -> names.contains(name));  
}
```

Every Method Must Return Something

- Unlike many imperative APIs, `void` is not an appropriate return type
- A reactive flow declares behavior but does not execute it
- Consumers must subscribe to a `Flux` or `Mono` in order to begin execution

Imperative Operation

```
void delete(String id) {  
    this.restTemplate.delete(URI, id);  
}
```

```
public static void main(String[] args) {  
    delete("test-id");  
}
```

Reactive Operation

```
Mono<Void> delete(String id) {  
    return this.httpClient.delete(URI, id);  
}  
  
public static void main(String[] args) {  
    CountdownLatch latch = new CountdownLatch(1);  
  
    delete("test-id")  
        .subscribe(n -> {}, Throwable::printStackTrace, () -> latch.countDown());  
  
    latch.await();  
}
```

Scope of a Method

- Methods should be designed to be small and reusable
- Reusable methods can be more easily composed into larger operations
- These methods can be more flexibly combined into parallel or sequential operations

Reusable Methods

```
Mono<ListApplicationsResponse> getPage(int page) {  
    return this.client.applicationsV2()  
        .list(ListApplicationsRequest.builder()  
            .page(page)  
            .build());  
}  
  
public static void main(String[] args) {  
    getPage(1)  
        .map(PaginatedResponse::getTotalPages)  
        .flatMap(totalPages -> Flux.range(2, totalPages)  
            .flatMap(page -> getPage(page)))  
        .subscribe(System.out::println);  
}
```

SpringOne Platform

Implement

Pivotal

Sequential and Parallel Coordination

- All significant performance improvements come from concurrency
- But, concurrency is very hard to do properly
- Reactive frameworks allow you to define sequential and parallel relationships between operations
- The framework does the hard work of coordinating the operations

Coordination

```
Mono<ListApplicationsResponse> getPage(int page) {  
    return this.client.applicationsV2()  
        .list(ListApplicationsRequest.builder()  
            .page(page)  
            .build());  
}  
  
public static void main(String[] args) {  
    getPage(1)  
        .map(PaginatedResponse::getTotalPages)  
        .flatMap(totalPages -> Flux.range(2, totalPages)  
            .flatMap(page -> getPage(page)))  
        .subscribe(System.out::println);  
}
```

Conditional Logic

- Like normal values, errors pass through the flow's operations
- These errors can be passed all the way to consumers, or flows can change behavior based on them
- Flows can transform errors or generate new results based on the error

Error Logic

```
public Mono<AppStatsResponse> getApp(GetAppRequest request) {  
    return client.applications()  
        .statistics(AppStatsRequest.builder()  
            .applicationId(request.id())  
            .build())  
        .otherwise(ExceptionUtils.statusCode(APP_STOPPED_ERROR),  
            t -> Mono.just(AppStatsResponse.builder().build()));  
}
```

Conditional Logic

- It is valid for a flow to complete without sending any elements
 - This is often equivalent to returning `null` from an imperative API
- This completion can be passed all the way to consumers
- Alternatively, flows can switch behavior or generate new elements based on the completion

Empty Logic

```
public Flux<GetDomainsResponse> getDomains(GetDomainsRequest request) {  
    return requestPrivateDomains(request.getId())  
        .switchIfEmpty(requestSharedDomains(request.getId()));  
}
```


Empty Logic

```
public Mono<String> getDomainId(GetDomainIdRequest request) {  
    return getPrivateDomainId(request.getName())  
        .otherwiseIfEmpty(getSharedDomainId(request.getName()))  
        .otherwiseIfEmpty(ExceptionUtils.illegalState(  
            "Domain %s not found", request.getName()));  
}
```

Conditional Logic

- Sometimes you want to make decisions not based on errors or emptiness, but on values themselves
- While it's possible to implement this logic using operators, it often proves to be more complex than is worthwhile
- It's OK to use imperative conditional statements

Conditional Logic Avoiding Imperative

```
public Mono<String> getDomainId(String domain, String organizationId) {  
    return Mono.just(domain)  
        .filter(d -> d == null)  
        .then(getSharedDomainIds())  
            .switchIfEmpty(getPrivateDomainIds(organizationId))  
            .next()  
            .otherwiseIfEmpty(ExceptionUtils.illegalState("Domain not found"))  
        .otherwiseIfEmpty(getPrivateDomainId(domain, organizationId))  
            .otherwiseIfEmpty(getSharedDomainId(domain))  
            .otherwiseIfEmpty(ExceptionUtils.illegalState("Domain %s not found", domain));  
}
```

Conditional Logic Using Imperative

```
public Mono<String> getDomainId(String domain, String organizationId) {  
    if (domain == null) {  
        return getSharedDomainIds()  
            .switchIfEmpty(getPrivateDomainIds(organizationId))  
            .next()  
            .otherwiseIfEmpty(ExceptionUtils.illegalState("Domain not found"));  
    } else {  
        return getPrivateDomainId(domain, organizationId)  
            .otherwiseIfEmpty(getSharedDomainId(domain))  
            .otherwiseIfEmpty(ExceptionUtils.illegalState("Domain %s not found", domain));  
    }  
}
```

Testing

- Most useful flows will be asynchronous
- Testing frameworks are aggressively synchronous, registering results before asynchronous results return
- To compensate for this you must block the main thread and make test assertions in the asynchronous thread


Test Finishing Before Assertions

```
@Test
public void noLatch() {
    Mono.just("alpha")
        .subscribeOn(Schedulers.computation())
        .subscribe(s -> assertEquals("bravo", s));
}
```

Test Finishing Before Assertions

@Test

```
public void noLatch() {  
    Mono.just("alpha")  
        .subscribeOn(Schedulers.computation())  
        .subscribe(s -> assertEquals("bravo", s));  
}
```



Test Blocking Using CountdownLatch

```
@Test
public void latch() throws InterruptedException {
    CountdownLatch latch = new CountdownLatch(1);
    AtomicReference<String> actual = new AtomicReference<>();

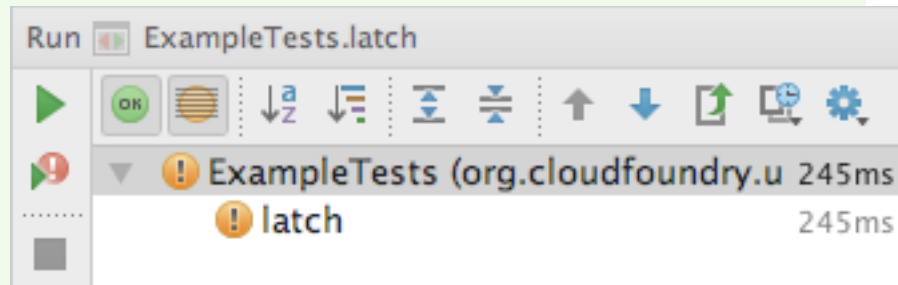
    Mono.just("alpha")
        .subscribeOn(Schedulers.computation())
        .subscribe(actual::set, t -> latch.countDown(), latch::countDown);

    latch.await();
    assertEquals("bravo", actual.get());
}
```


Test Blocking Using CountDownLatch

@Test

```
public void latch() throws InterruptedException {  
    CountDownLatch latch = new CountDownLatch(1);  
    AtomicReference<String> actual = new AtomicReference<>();  
  
    Mono.just("alpha")  
        .subscribeOn(Schedulers.computation())  
        .subscribe(actual::set, t -> latch.countDown(), latch::countDown);  
  
    latch.await();  
    assertEquals("bravo", actual.get());  
}
```



Test Subscriber

- Testing a Reactive design requires more than just blocking
 - Multiple value assertion
 - Expected error assertion
 - Unexpected error failure
- A `TestSubscriber<T>` can collect this behavior into a single type

Test Using TestSubscriber<T>

```
@Test
public void testSubscriber() throws InterruptedException {
    TestSubscriber<String> testSubscriber = new TestSubscriber<>();

    Flux.just("alpha", "bravo")
        .subscribeOn(Schedulers.computation())
        .subscribe(testSubscriber
            .assertEquals("alpha")
            .assertEquals("bravo"));

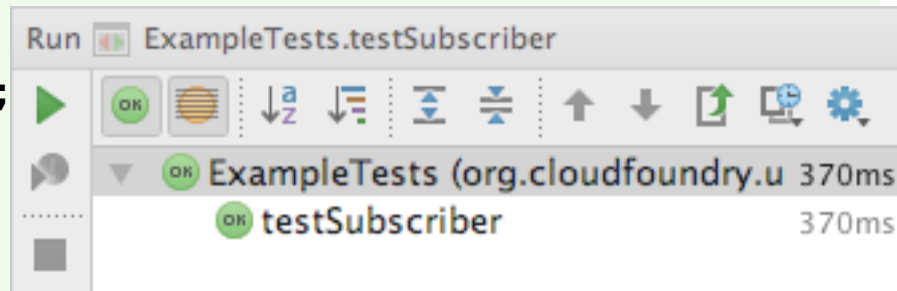
    testSubscriber.verify(ofSeconds(5));
}
```

Test Using TestSubscriber<T>

```
@Test
public void testSubscriber() throws InterruptedException {
    TestSubscriber<String> testSubscriber = new TestSubscriber<>();

    Flux.just("alpha", "bravo")
        .subscribeOn(Schedulers.computation())
        .subscribe(testSubscriber
            .assertEquals("alpha")
            .assertEquals("bravo"));

    testSubscriber.verify(ofSeconds(5));
}
```



SpringOne Platform

Use

Pivotal

Countdown Latches

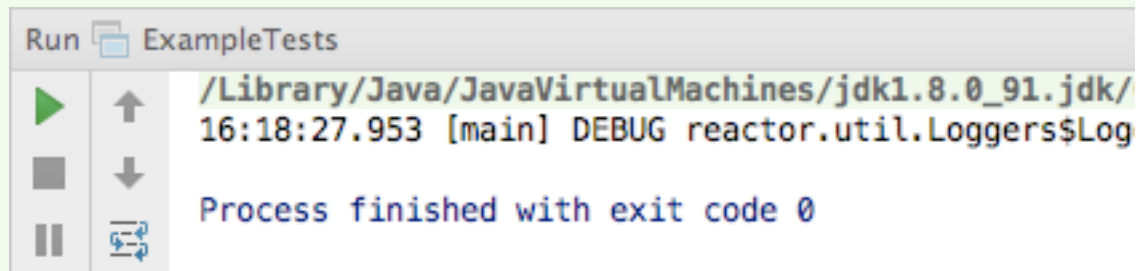
- Reactive frameworks can only coordinate their own operations
- Many execution environments have processes that outlast individual threads
 - Reactive programming plays nicely here
- Some execution environments aggressively terminate the process before any individual thread

main() Finishing Before Subscriber<T>

```
public static void main(String[] args) {  
    Mono.just("alpha")  
        .delaySubscription(Duration.ofSeconds(1))  
        .subscribeOn(Schedulers.computation())  
        .subscribe(System.out::println);  
}
```

main() Finishing Before Subscriber<T>

```
public static void main(String[] args) {  
    Mono.just("alpha")  
        .delaySubscription(Duration.ofSeconds(1))  
        .subscribeOn(Schedulers.computation())  
        .subscribe(System.out::println);  
}
```



```
Run ExampleTests  
/Library/Java/JavaVirtualMachines/jdk1.8.0_91.jdk/  
16:18:27.953 [main] DEBUG reactor.util.Loggers$Log  
  
Process finished with exit code 0
```


main() Waiting for Subscriber<T>

```
public static void main(String[] args) throws InterruptedException {  
    CountdownLatch latch = new CountdownLatch(1);  
  
    Mono.just("alpha")  
        .delaySubscription(Duration.ofSeconds(1))  
        .subscribeOn(Schedulers.computation())  
        .subscribe(System.out::println, t -> latch.countDown(),  
                    latch::countDown);  
  
    latch.await();  
}
```

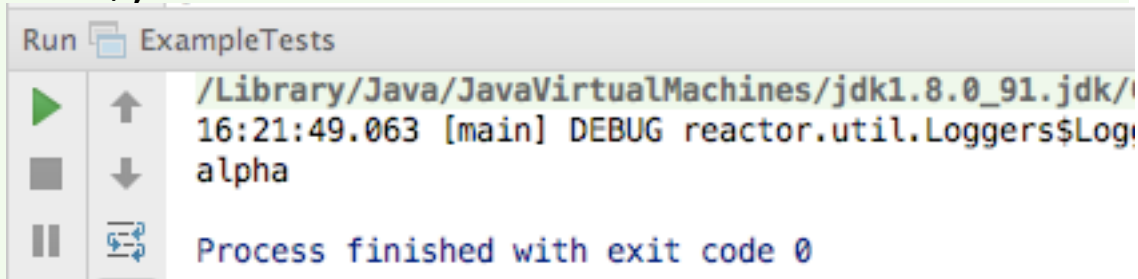
main() Waiting for Subscriber<T>

```
public static void main(String[] args) throws InterruptedException {  
    CountdownLatch latch = new CountdownLatch(1);
```

```
    Mono.just("alpha")  
        .delaySubscription(Duration.ofSeconds(1))  
        .subscribeOn(Schedulers.computation())  
        .subscribe(System.out::println, t -> latch.countDown(),  
                    latch::countDown);
```

```
    latch.await();
```

```
}
```



Blocking flows

- It's very common to bridge between imperative and reactive APIs
- In those cases blocking while waiting for a result is appropriate.

Bridging Imperative and Reactive using `block()`

```
Mono<User> requestUser(String name) {...}
```

```
@Override
```

```
    User requestUser(String name) {  
        return getUser(name)  
            .block();  
    }
```

Bridging Imperative and Reactive using `block()`

```
Flux<User> requestUsers() {...}
```

```
@Override
```

```
List<User> listUsers() {  
    return requestUsers()  
        .collectList()  
        .block();  
}
```

Error Handling

- Because errors are values they must be handled, they do not just show up
- Subscribe has 0...3 parameters
- Release latches on error

Error Handling

```
public static void main(String[] args) throws InterruptedException {  
    CountdownLatch latch = new CountdownLatch(1);  
  
    Flux.concat(Mono.just("alpha"), Mono.error(new IllegalStateException()))  
        .subscribe(System.out::println, t -> {  
            t.printStackTrace();  
            latch.countDown();  
        }, latch::countDown);  
  
    latch.await();  
}
```

Composable Method References

- Reactive programming is susceptible to callback hell, but it doesn't have to be
- Extracting behavior into private methods can be even more valuable than in imperative programming
- Well named methods, and method reference syntax, lead to very readable flows

Composable Method References

```
@Override
public Flux<ApplicationSummary> list() {
    return Mono
        .when(this.cloudFoundryClient, this.spaceId)
        .then(function(DefaultApplications::requestSpaceSummary))
        .flatMap(DefaultApplications::extractApplications)
        .map(DefaultApplications::toApplicationSummary);
}
```

Point Free

- You may have noticed that our code samples use a *very* compact style
 - Point Free (<http://bit.ly/Point-Free>)
- It helps the developer think about composing functions (high level), rather than shuffling data (low level)
- You don't have to use, but we find that most people prefer it (eventually)

Point Free Style

```
Mono<Void> deleteApplication(String name) {  
    return PaginationUtils  
        .requestClientV2Resources(page -> this.client.applicationsV2()  
            .list(ListApplicationsRequest.builder()  
                .name(name)  
                .page(page)  
                .build()))  
        .single()  
        .map(applicationResource -> applicationResource.getMetadata().getId())  
        .then(applicationId -> this.client.applicationsV2()  
            .delete>DeleteApplicationRequest.builder()  
                .applicationId(applicationId)  
                .build())));  
}
```

SpringOne Platform

Summary

Pivotal

Designing, Implementing and Using Reactive APIs

- Design
 - When to use
 - Return Mono/Flux
 - Reusable Methods
- Implement
 - Sequential and Parallel
 - Conditional logic
 - Testing
- Use
 - Countdown Latches
 - Blocking flows
 - Error Handling
 - Composable methods
 - Point Free Style

SpringOne Platform

Learn More. Stay Connected.

<https://github.com/cloudfoundry/cf-java-client>



@springcentral
spring.io/blog



@pivotal
pivotal.io/blog



@pivotalcf
<http://engineering.pivotal.io>